


Data Contract Specification

Stars 318 Slack join chat



Data contracts bring data providers and data consumers together.

A *data contract* is a document that defines the structure, format, semantics, quality, and terms of use for exchanging data between a data provider and their consumers. Think of an API, but for data. A data contract is implemented by a data product or other data technologies, even legacy data warehouses. Data contracts can also be used for the input port to specify the expectations of data dependencies and verify given guarantees.

The *data contract specification* defines a YAML format to describe attributes of provided data sets. It is data platform neutral and can be used with any data platform, such as AWS S3, Google BigQuery, Azure, Databricks, and Snowflake. The data contract specification is an open initiative to define a common data contract format. It follows [OpenAPI](#) and [AsyncAPI](#) conventions.

Data contracts come into play when data is exchanged between different teams or organizational units, such as in a [data mesh architecture](#). First, and foremost, data contracts are a communication tool to express a common understanding of how data should be structured and interpreted. They make semantic and quality expectations explicit. They are often created collaboratively in [workshops](#) together with data providers and data consumers. Later in development and production, they also serve as the basis for code generation, testing, schema validations, quality checks, monitoring, access control, and computational governance policies.

The specification comes along with the [Data Contract CLI](#), an open-source tool to develop, validate, and enforce data contracts.

Note: The term "data contract" refers to a specification that is usually owned by the data provider and thus does not align with a "contract" in a legal sense as a mutual agreement

between two parties. The term “contract” may be somewhat misleading, but it is how it is used by the industry. The mutual agreement between one data provider and one data consumer is the “data usage agreement” that refers to a data contract. Data usage agreements have a defined lifecycle, start/end date, and help the data provider to track who accesses their data and for which purposes.

Version

1.1.0([Changelog](#))

Example

View in [Data Contract Catalog](#)

```
dataContractSpecification: 1.1.0
id: urn:datacontract:checkout:orders-latest
info:
  title: Orders Latest
  version: 2.0.0
  description: |
    Successful customer orders in the webshop.
    All orders since 2020-01-01.
    Orders with their line items are in their current state (no history included).
  owner: Checkout Team
  contact:
    name: John Doe (Data Product Owner)
    url: https://teams.microsoft.com/l/channel/example/checkout
servers:
  production:
    type: s3
    environment: prod
    location: s3://datacontract-example-orders-latest/v2/{model}/*.json
    format: json
    delimiter: new_line
    description: "One folder per model. One file per day."
    roles:
      - name: analyst_us
        description: Access to the data for US region
      - name: analyst_cn
        description: Access to the data for China region
terms:
  usage: |
    Data can be used for reports, analytics and machine learning use cases.
    Order may be linked and joined by other tables
  limitations: |
    Not suitable for real-time use cases.
    Data may not be used to identify individual customers.
    Max data processing per day: 10 TiB
policies:
```

```
- name: privacy-policy
  url: https://example.com/privacy-policy
- name: license
  description: External data is licensed under agreement 1234.
  url: https://example.com/license/1234
billing: 5000 USD per month
noticePeriod: P3M
models:
orders:
  description: One record per order. Includes cancelled and deleted orders.
  type: table
  fields:
    order_id:
      $ref: '#/definitions/order_id'
      required: true
      unique: true
      primaryKey: true
    order_timestamp:
      description: The business timestamp in UTC when the order was successfully
      type: timestamp
      required: true
      examples:
        - "2024-09-09T08:30:00Z"
      tags: ["business-timestamp"]
    order_total:
      description: Total amount the smallest monetary unit (e.g., cents).
      type: long
      required: true
      examples:
        - 9999
    quality:
      - type: sql
        description: 95% of all order total values are expected to be between
        query: |
          SELECT quantile_cont(order_total, 0.95) AS percentile_95
          FROM orders
        mustBeBetween: [1000, 49900]
  customer_id:
    description: Unique identifier for the customer.
    type: text
    minLength: 10
    maxLength: 20
  customer_email_address:
    description: The email address, as entered by the customer.
    type: text
    format: email
    required: true
    pii: true
    classification: sensitive
    quality:
      - type: text
        description: The email address is not verified and may be invalid.
  lineage:
```

```

    inputFields:
      - namespace: com.example.service.checkout
        name: checkout_db.orders
        field: email_address
    processed_timestamp:
      description: The timestamp when the record was processed by the data platf
      type: timestamp
      required: true
      config:
        jsonType: string
        jsonFormat: date-time
    quality:
      - type: sql
        description: The maximum duration between two orders should be less that 3
        query: |
          SELECT MAX(duration) AS max_duration
          FROM (
            SELECT EXTRACT(EPOCH FROM (order_timestamp - LAG(order_timestamp) OVER
              FROM orders
            )
          )
        mustBeLessThan: 3600
      - type: sql
        description: Row Count
        query: |
          SELECT count(*) as row_count
          FROM orders
        mustBeGreaterThan: 5
    examples:
      - |
        order_id,order_timestamp,order_total,customer_id,customer_email_address,pr
        "1001","2030-09-09T08:30:00Z",2500,"1000000001","mary.taylor82@example.com
        "1002","2030-09-08T15:45:00Z",1800,"1000000002","michael.miller83@example.
        "1003","2030-09-07T12:15:00Z",3200,"1000000003","michael.smith5@example.co
        "1004","2030-09-06T19:20:00Z",1500,"1000000004","elizabeth.moore80@example
        "1005","2030-09-05T10:10:00Z",4200,"1000000004","elizabeth.moore80@example
        "1006","2030-09-04T14:55:00Z",2800,"1000000005","john.davis28@example.com"
        "1007","2030-09-03T21:05:00Z",1900,"1000000006","linda.brown67@example.com
        "1008","2030-09-02T17:40:00Z",3600,"1000000007","patricia.smith40@example.
        "1009","2030-09-01T09:25:00Z",3100,"1000000008","linda.wilson43@example.co
        "1010","2030-08-31T22:50:00Z",2700,"1000000009","mary.smith98@example.com"
    line_items:
      description: A single article that is part of an order.
      type: table
      fields:
        line_item_id:
          type: text
          description: Primary key of the lines_item_id table
          required: true
        order_id:
          $ref: '#/definitions/order_id'
          references: orders.order_id
      sku:
        description: The purchased article number

```

```
    $ref: '#/definitions/sku'
primaryKey: ["order_id", "line_item_id"]
examples:
  - |
    line_item_id,order_id,sku
    "LI-1","1001","5901234123457"
    "LI-2","1001","4001234567890"
    "LI-3","1002","5901234123457"
    "LI-4","1002","2001234567893"
    "LI-5","1003","4001234567890"
    "LI-6","1003","5001234567892"
    "LI-7","1004","5901234123457"
    "LI-8","1005","2001234567893"
    "LI-9","1005","5001234567892"
    "LI-10","1005","6001234567891"
```

definitions:

```
order_id:
  title: Order ID
  type: text
  format: uuid
  description: An internal ID that identifies an order in the online shop.
  examples:
    - 243c25e5-a081-43a9-aeab-6d5d5b6cb5e2
  pii: true
  classification: restricted
  tags:
    - orders
```

```
sku:
  title: Stock Keeping Unit
  type: text
  pattern: ^[A-Za-z0-9]{8,14}$
  examples:
    - "96385074"
  description: |
    A Stock Keeping Unit (SKU) is an internal unique identifier for an article.
    It is typically associated with an article's barcode, such as the EAN/GTIN.
  links:
    wikipedia: https://en.wikipedia.org/wiki/Stock\_keeping\_unit
  tags:
    - inventory
```

servicelevels:

```
availability:
  description: The server is available during support hours
  percentage: 99.9%
retention:
  description: Data is retained for one year
  period: P1Y
  unlimited: false
```

```
latency:
  description: Data is available within 25 hours after the order was placed
  threshold: 25h
sourceTimestampField: orders.order_timestamp
processedTimestampField: orders.processed_timestamp
```


```
freshness:
  description: The age of the youngest row in a table.
  threshold: 25h
  timestampField: orders.order_timestamp
frequency:
  description: Data is delivered once a day
  type: batch # or streaming
  interval: daily # for batch, either or cron
  cron: 0 0 * * * # for batch, either or interval
support:
  description: The data is available during typical business hours at headquarte
  time: 9am to 5pm in EST on business days
  responseTime: 1h
backup:
  description: Data is backed up once a week, every Sunday at 0:00 UTC.
  interval: weekly
  cron: 0 0 * * 0
  recoveryTime: 24 hours
  recoveryPoint: 1 week
tags:
  - checkout
  - orders
  - s3
links:
  datacontractCli: https://cli.datacontract.com
```

Data Contract CLI

The [Data Contract CLI](#) is a command line tool and Python library to lint, test, import and export data contracts.

Here is short example how to verify that your actual dataset matches the data contract:

```
pip3 install "datacontract-cli[all]"
datacontract test https://datacontract.com/examples/orders-latest/datacontract.yam
```



or, if you prefer Docker:

```
docker run datacontract/cli test https://datacontract.com/examples/orders-latest/d
```



The Data Contract contains all required information to verify data:


- The *servers* block has the connection details to the actual data set.
- The *models* define the syntax, formats, and constraints.

- The *quality* defined further quality checks.

The Data Contract CLI chooses the appropriate engine, formulates test cases, connects to the server, and executes the tests, based on the server type.

More information and configuration options on cli.datacontract.com.

Specification



- [Data Contract Object](#)
- [Info Object](#)
- [Contact Object](#)
- [Server Object](#)
- [Terms Object](#)
- [Model Object](#)
- [Field Object](#)
- [Definition Object](#)
- [Service Level Object](#)
- [Quality Object](#)
- [Lineage Object](#)
- [Data Types](#)
- [Specification Extensions](#)

[JSON Schema](#) of the Data Contract Specification.

Data Contract Object

This is the root document.

It is *RECOMMENDED* that the root document be named: `datacontract.yaml` .

Field	Type	Description
dataContractSpecification	string	REQUIRED. Specifies the Data Contract Specification being used.
id	string	REQUIRED. An organization-wide unique technical identifier, such as a UUID, URN, slug, string, or number
info	Info Object	REQUIRED. Specifies the metadata of the data contract. May be used by tooling.
servers	Map[string , Server Object]	Specifies the servers of the data contract.
terms	Terms Object	Specifies the terms and conditions of the data contract.
models	Map[string , Model Object]	Specifies the logical data model.
definitions	Map[string , Definition Object]	Specifies definitions.
servicelevels	Service Levels Object	Specifies the service level of the provided data
links	Map[string , string]	Additional external documentation links.
tags	Array of string	Custom metadata to provide additional context.

This object *MAY* be extended with [Specification Extensions](#).

Info Object

Metadata and life cycle information about the data contract.

Field	Type	Description
title	string	REQUIRED. The title of the data contract.

Field	Type	Description
version	string	REQUIRED. The version of the data contract document (which is distinct from the Data Contract Specification version or the Data Product implementation version).
status	string	The status of the data contract. Can be <code>proposed</code> , <code>in development</code> , <code>active</code> , <code>deprecated</code> , <code>retired</code> .
description	string	A description of the data contract.
owner	string	The owner or team responsible for managing the data contract and providing the data.
contact	Contact Object	Contact information for the data contract.

This object *MAY* be extended with [Specification Extensions](#).

Contact Object

Contact information for the data contract.

Field	Type	Description
name	string	The identifying name of the contact person/organization.
url	string	The URL pointing to the contact information. This <i>MUST</i> be in the form of a URL.
email	string	The email address of the contact person/organization. This <i>MUST</i> be in the form of an email address.

This object *MAY* be extended with [Specification Extensions](#).

Server Object

The fields are dependent on the defined type.

Field	Type	Description
type	string	REQUIRED. The type of the data product technology that implements the data contract. Well-known server types are: <code>bigquery</code> , <code>s3</code> , <code>glue</code> , <code>redshift</code> , <code>azure</code> , <code>sqlserver</code> , <code>snowflake</code> , <code>databricks</code> , <code>postgres</code> , <code>oracle</code> , <code>kafka</code> , <code>pubsub</code> , <code>sftp</code> , <code>kinesis</code> , <code>trino</code> , <code>local</code>

Field	Type	Description
description	string	An optional string describing the server.
environment	string	An optional string describing the environment, e.g., prod, sit, stg.
roles	Array of Server Role Object	An optional array of roles that are available and can be requested to access the server for role-based access control. E.g. separate roles for different regions or sensitive data.

This object *MAY* be extended with [Specification Extensions](#).

BigQuery Server Object

Field	Type	Description
type	string	bigquery
project	string	The GCP project name.
dataset	string	

S3 Server Object

Field	Type	Description
type	string	s3
location	string	S3 URL, starting with s3://
endpointUrl	string	The server endpoint for S3-compatible servers, such as MioIO or Google Cloud Storage, e.g., https://minio.example.com
format	string	Format of files, such as parquet , delta , json , csv
delimiter	string	(Only for format = json), how multiple json documents are delimited within one file, e.g., new_line , array

Example (AWS S3):

```
servers:
  production:
    type: s3
    location: s3://acme-orders-prod/orders/
    format: json
    delimiter: new_line
```

Example (MinIO):

```
servers:
  minio:
    type: s3
    endpointUrl: http://localhost:9000
    location: s3://my-bucket/path/
    format: delta
```

Example (Google Cloud Storage):

```
servers:
  gcs:
    type: s3
    endpointUrl: https://storage.googleapis.com
    location: s3://my-bucket/path/**/*.parquet
    format: parquet
```

Redshift Server Object

Field	Type	Description
type	string	redshift
account	string	
database	string	
schema	string	
clusterIdentifier	string	Identifier of the cluster. Example: analytics-cluster
host	string	Host of the cluster. Example: analytics-cluster.example.eu-west-1.redshift.amazonaws.com
port	number	Port of the cluster. Example: 5439
endpoint	string	Endpoint of the cluster Example: analytics-cluster.example.eu-west-1.redshift.amazonaws.com:5439/analytics

Example, specifying an endpoint:

```
servers:
  analytics:
```

```
type: redshift
account: '123456789012'
database: analytics
schema: analytics
endpoint: analytics-cluster.example.eu-west-1.redshift.amazonaws.com:5439/anal
```

Example, specifying the cluster identifier:

```
servers:
  analytics:
    type: redshift
    account: '123456789012'
    database: analytics
    schema: analytics
    clusterIdentifier: analytics-cluster
```

Example, specifying the cluster host:

```
servers:
  analytics:
    type: redshift
    account: '123456789012'
    database: analytics
    schema: analytics
    host: analytics-cluster.example.eu-west-1.redshift.amazonaws.com
    port: 5439
```

Azure Server Object

Field	Type	Description
type	string	azure
storageAccount	string	The storage account name that contains the files
location	string	Path to Azure Blob Storage or Azure Data Lake Storage (ADLS) supports globs. Starting with az:// or abfss Recommended pattern is abfss://<container_name>/<path> az://my_storage_account_name.blob.core.windows.net/my_ or abfss://my_container_name/path/*.parquet
format	string	Format of files, such as parquet , json , csv
delimiter	string	(Only for format = json), how multiple json documents are delimited new_line , array

SQL-Server Server Object

Field	Type	Description
type	string	sqlserver
host	string	The host to the database server
port	integer	The port to the database server, default: 1433
database	string	The name of the database, e.g., database .
schema	string	The name of the schema in the database, e.g., dbo .
driver	string	The name of the supported driver, e.g., ODBC Driver 18 for SQL Server .

Snowflake Server Object

Field	Type	Description
type	string	snowflake
account	string	
database	string	
schema	string	

Databricks Server Object

Field	Type	Description
type	string	databricks
host	string	The Databricks host, e.g., dbc-abcdefgh-1234.cloud.databricks.com
catalog	string	The name of the Hive or Unity catalog
schema	string	The schema name in the catalog

Postgres Server Object

Field	Type	Description
type	string	postgres
host	string	The host to the database server
port	integer	The port to the database server
database	string	The name of the database, e.g., postgres .

Field	Type	Description
schema	string	The name of the schema in the database, e.g., <code>public</code> .

Oracle Server Object

Field	Type	Description
type	string	oracle
host	string	The host to the oracle server
port	integer	The port to the oracle server
serviceName	string	The name of the service

Kafka Server Object

Field	Type	Description
type	string	kafka
host	string	The bootstrap server of the kafka cluster.
topic	string	The topic name.
format	string	The format of the message. Examples: json, avro, protobuf. Default: json.

Pub/Sub Server Object

Field	Type	Description
type	string	pubsub
project	string	The GCP project name.
topic	string	The topic name.

sftp Server Object

Field	Type	Description
type	string	sftp
location	string	S3 URL, starting with <code>sftp://</code>
format	string	Format of files, such as <code>parquet</code> , <code>delta</code> , <code>json</code> , <code>csv</code>

Field	Type	Description
delimiter	string	(Only for format = json), how multiple json documents are delimited within one file, e.g., new_line , array

AWS Kinesis Data Streams Server Object

Field	Type	Description
type	string	kinesis
stream	string	The name of the Kinesis data stream.
region	string	AWS region, e.g., eu-west-1 .
format	string	The format of the records. Examples: json, avro, protobuf.

Trino Server Object

Field	Type	Description
type	string	trino
host	string	The Trino host
port	integer	The Trino port
catalog	string	The name of the catalog, e.g., my_catalog .
schema	string	The name of the schema in the catalog, e.g., my_schema .

Local Server Object

Field	Type	Description
type	string	local
path	string	The relative or absolute path to the data file(s), such as ./folder/data.parquet .
format	string	The format of the file(s), such as parquet , delta , csv , or json .

Server Role Object

Field	Type	Description
name	string	Name of the role
description	string	A description of the role and what access the role provides.

Terms Object

The terms and conditions of the data contract.

Field	Type	Description
usage	string	The usage describes the way the data is expected to be used. Can contain business and technical information.
limitations	string	The limitations describe the restrictions on how the data can be used, can be technical or restrictions on what the data may not be used for.
policies	Array of Policy Object	A list of policies, licenses, standards, that are applicable for this data contract and that must be acknowledged by data consumers.
billing	string	The billing describes the pricing model for using the data, such as whether it's free, having a monthly fee, or metered pay-per-use.
noticePeriod	string	The period of time that must be given by either party to terminate or modify a data usage agreement. Uses ISO-8601 period format, e.g., P3M for a period of three months.

This object *MAY* be extended with [Specification Extensions](#).

Policy Object

Field	Type	Description
name	string	Name of the policy.
description	string	A description of the policy.
url	string	An URL that refers to the policy.

Model Object

The Model Object describes the structure and semantics of a data model, such as tables, views, or structured files.

The name of the data model (table name) is defined by the key that refers to this Model Object.

Field	Type	Description
type	string	The type of the model. Examples: <code>table</code> , <code>view</code> , <code>object</code> . Default: <code>table</code> .
description	string	An optional string describing the data model.
title	string	An optional string for the title of the data model. Especially useful if the name of the model is cryptic or contains abbreviations.
fields	Map[string , Field Object]	The fields (e.g. columns) of the data model.
primaryKey	Array of string	If the primary key is a compound key, list the field names that constitute the primary key. Alternative to field-level <code>primaryKey</code> .
quality	Array of Quality Object	Specifies the quality attributes on model level.
examples	Array of Any	Specifies example data sets for the model.
config	Config Object	Any additional key-value pairs that might be useful for further tooling.

This object *MAY* be extended with [Specification Extensions](#).

Field Object

The Field Objects describes one field (column, property, nested field) of a data model.

Field	Type	Description
description	string	An optional string describing the semantic of the data in this field.
type	Data Type	The logical data type of the field.
title	string	An optional string providing a human readable name for the field. Especially useful if the field name is cryptic or contains abbreviations.
enum	array of string	A value must be equal to one of the elements in this array value. Only evaluated if the value is not null.

Field	Type	Description
required	boolean	An indication, if this field must contain a value and may not be null. Default: false
primaryKey	boolean	If this field is a primary key. Default: false
references	string	The reference to a field in another model. E.g. use 'orders.order_id' to reference the order_id field of the model orders. Think of defining a foreign key relationship.
unique	boolean	An indication, if the value must be unique within the model. Default: false
format	string	<p>email : A value must be complaint to RFC 5321, section 4.1.2.</p> <p>uri : A value must be complaint to RFC 3986.</p> <p>uuid : A value must be complaint to RFC 4122. Only evaluated if the value is not null. Only applies to unicode character sequences types (string , text , varchar).</p>
precision	number	The maximum number of digits in a number. Only applies to numeric values. Defaults to 38.
scale	number	The maximum number of decimal places in a number. Only applies to numeric values. Defaults to 0.
minLength	number	A value must greater than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to unicode character sequences types (string , text , varchar).
maxLength	number	A value must less than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to unicode character sequences types (string , text , varchar).

Field	Type	Description
pattern	string	A value must be valid according to the ECMA-262 regular expression dialect. Only evaluated if the value is not null. Only applies to unicode character sequences types (<code>string</code> , <code>text</code> , <code>varchar</code>).
minimum	number	A value of a number must greater than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to numeric values.
exclusiveMinimum	number	A value of a number must greater than the value of this. Only evaluated if the value is not null. Only applies to numeric values.
maximum	number	A value of a number must less than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to numeric values.
exclusiveMaximum	number	A value of a number must less than the value of this. Only evaluated if the value is not null. Only applies to numeric values.
example	string	DEPRECATED, use examples. An example value.
examples	Array of Any	A list of example values.
pii	boolean	An indication, if this field contains Personal Identifiable Information (PII).
classification	string	The data class defining the sensitivity level for this field, according to the organization's classification scheme. Examples may be: <code>sensitive</code> , <code>restricted</code> , <code>internal</code> , <code>public</code> .
tags	Array of string	Custom metadata to provide additional context.
links	Map[string , string]	Additional external documentation links.

Field	Type	Description
\$ref	string	A reference URI to a definition in the specification, internally or externally. Properties will be inherited from the definition.
fields	Map[string , Field Object]	The nested fields (e.g. columns) of the object, record, or struct. Use only when type is <code>object</code> , <code>record</code> , or <code>struct</code> .
items	Field Object	The type of the elements in the array. Use only when type is <code>array</code> .
keys	Field Object	Describes the key structure of a map. Defaults to <code>type: string</code> if a map is defined as type. Not all server types support different key types. Use only when type is <code>map</code> .
values	Field Object	Describes the value structure of a map. Use only when type is <code>map</code> .
quality	Array of Quality Object	Specifies the quality attributes on field level.
lineage	Lineage Object	Provides information where the data comes from.
config	Config Object	Any additional key-value pairs that might be useful for further tooling.

This object *MAY* be extended with [Specification Extensions](#).

Definition Object

The Definition Object includes a clear and concise explanations of syntax, semantic, and classification of a business object in a given domain. It serves as a reference for a common understanding of terminology, ensure consistent usage and to identify join-able fields. Models fields can refer to definitions using the `$ref` field to link to existing definitions and avoid duplicate documentations.

Field	Type	Description
type	Data Type	REQUIRED. The logical data type
title	string	The business name of this definition.

Field	Type	Description
description	string	Clear and concise explanations related to the domain
enum	array of string	A value must be equal to one of the elements in this array value. Only evaluated if the value is not null.
format	string	<p>email : A value must be complaint to RFC 5321, section 4.1.2.</p> <p>uri : A value must be complaint to RFC 3986.</p> <p>uuid : A value must be complaint to RFC 4122.</p> <p>Only evaluated if the value is not null. Only applies to unicode character sequences types (string , text , varchar).</p>
precision	number	The maximum number of digits in a number. Only applies to numeric values. Defaults to 38.
scale	number	The maximum number of decimal places in a number. Only applies to numeric values. Defaults to 0.
minLength	number	A value must greater than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to unicode character sequences types (string , text , varchar).
maxLength	number	A value must less than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to unicode character sequences types (string , text , varchar).
pattern	string	A value must be valid according to the ECMA-262 regular expression dialect. Only evaluated if the value is not null. Only applies to unicode character sequences types (string , text , varchar).
minimum	number	A value of a number must greater than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to numeric values.
exclusiveMinimum	number	A value of a number must greater than the value of this. Only evaluated if the value is not null. Only applies to numeric values.

Field	Type	Description
maximum	number	A value of a number must less than, or equal to, the value of this. Only evaluated if the value is not null. Only applies to numeric values.
exclusiveMaximum	number	A value of a number must less than the value of this. Only evaluated if the value is not null. Only applies to numeric values.
examples	Array of Any	A list of example values.
pii	boolean	An indication, if this field contains Personal Identifiable Information (PII).
classification	string	The data class defining the sensitivity level for this field, according to the organization's classification scheme.
tags	Array of string	Custom metadata to provide additional context.
links	Map[string , string]	Additional external documentation links.
fields	Map[string , Field Object]	The nested fields (e.g. columns) of the object, record, or struct. Use only when type is object , record , or struct .
items	Field Object	The type of the elements in the array. Use only when type is array .
keys	Field Object	Describes the key structure of a map. Defaults to type: string if a map is defined as type. Not all server types support different key types. Use only when type is map .
values	Field Object	Describes the value structure of a map. Use only when type is map .

This object *MAY* be extended with [Specification Extensions](#).

Service Levels Object

A service level is defined as an agreed-upon, measurable level of performance for provided the data. Data Contract Specification defines well-known service levels. This list can be extended with custom service levels.

One can either describe each service level informally using the `description` field, or make use of the predefined fields for automation support, e.g., via the [Data Contract CLI](#).

Field	Type	Description
availability	Availability Object	The promised uptime of the system that provides the data
retention	Retention Object	The period how long data will be available.
latency	Latency Object	The maximum amount of time from the source to its destination.
freshness	Freshness Object	The maximum age of the youngest entry.
frequency	Frequency Object	The update frequency.
support	Support Object	The times when support is provided.
backup	Backup Object	The details about data backup procedures.

This object *MAY* be extended with [Specification Extensions](#).

Availability Object

Availability refers to the promise or guarantee by the service provider about the uptime of the system that provides the data.

Field	Type	Description
description	string	An optional string describing the availability service level.
percentage	string	An optional string describing the guaranteed uptime in percent (e.g., 99.9%)

This object *MAY* be extended with [Specification Extensions](#).

Retention Object

Retention covers the period how long data will be available.

Field	Type	Description
description	string	An optional string describing the retention service level.

Field	Type	Description
period	string	An optional period of time, how long data is available. Supported formats: Simple duration (e.g., 1 year , 30d) and ISO 8601 duration (e.g, P1Y).
unlimited	boolean	An optional indicator that data is kept forever.
timestampField	string	An optional reference to the field that contains the timestamp that the period refers to.

This object *MAY* be extended with [Specification Extensions](#).

Latency Object

Latency refers to the maximum amount of time from the source to its destination.

Examples are the maximum duration it takes after an order has been recorded in the ecommerce shop until it is available in the orders table in the data analytics platform. This includes the waiting times until the next batch run is started and the processing time of the pipeline.

Field	Type	Description
description	string	An optional string describing the latency service level.
threshold	string	An optional maximum duration between the source timestamp and the processed timestamp. Supported formats: Simple duration (e.g., 24 hours , 5s) and ISO 8601 duration (e.g, PT24H).
sourceTimestampField	string	An optional reference to the field that contains the timestamp when the data was provided at the source.
processedTimestampField	string	An optional reference to the field that contains the processing timestamp, which denotes when the data is made available to consumers of this data contract.

This object *MAY* be extended with [Specification Extensions](#).

Freshness Object

Freshness refers to the maximum age of the youngest entry.

Field	Type	Description
description	string	An optional string describing the freshness service level.
threshold	string	An optional maximum age of the youngest entry. Supported formats: Simple duration (e.g., 24 hours , 5s) and ISO 8601 duration (e.g, PT24H).
timestampField	string	An optional reference to the field that contains the timestamp that the threshold refers to.

This object *MAY* be extended with [Specification Extensions](#).

Frequency Object

Frequency describes how often data is updated.

Field	Type	Description
description	string	An optional string describing the frequency service level.
type	string	An optional type of data processing. Typical values are batch , micro-batching , streaming , manual .
interval	string	Optional. Only for batch: How often the pipeline is triggered, e.g., daily .
cron	string	Optional. Only for batch: A cron expression when the pipelines is triggered. E.g., 0 0 * * * .

This object *MAY* be extended with [Specification Extensions](#).

Support Object

Support describes the times when support will be available for contact.

Field	Type	Description
description	string	An optional string describing the support service level.
time	string	An optional string describing the times when support will be available for contact such as 24/7 or business hours only .
responseTime	string	An optional string describing the time it takes for the support team to acknowledge a request. This does not mean the issue will be resolved immediately, but it assures users that their request has been received and will be dealt with.

This object *MAY* be extended with [Specification Extensions](#).

Backup Object

Backup specifies details about data backup procedures.

Field	Type	Description
description	string	An optional string describing the backup service level.
interval	string	An optional interval that defines how often data will be backed up, e.g., <code>daily</code> .
cron	string	An optional cron expression when data will be backed up, e.g., <code>0 0 * * *</code> .
recoveryTime	string	An optional Recovery Time Objective (RTO) specifies the maximum amount of time allowed to restore data from a backup after a failure or loss event (e.g., 4 hours, 24 hours).
recoveryPoint	string	An optional Recovery Point Objective (RPO) defines the maximum acceptable age of files that must be recovered from backup storage for normal operations to resume after a disaster or data loss event. This essentially measures how much data you can afford to lose, measured in time (e.g., 4 hours, 24 hours).

Quality Object

The quality object defines quality attributes.

Quality attributes are checks that can be applied to the data to ensure its quality. Data can be verified by executing these checks through a data quality engine.

Quality attributes can be:

- A text in natural language that describes the quality of the data.
- An individual SQL query that returns a single value that can be compared.
- Engine-specific types: Pre-defined quality checks, as defined by data quality libraries. Currently, the engines `soda` and `great-expectations` are supported.

A quality object can be specified on field level and on model level. The top-level quality object is deprecated.

Description Text

A description in natural language that defines the expected quality of the data. This is useful to express requirements or expectation when discussing the data contract with stakeholders.

Later in the development process, these might be translated into an executable check (such as `sql`). It can also be used as a prompt to check the data with an AI engine.

Field	Type	Description
type	string	text
description	string	A plain text describing the quality attribute in natural language.

Example:

```
models:
  my_table:
    fields:
      account_iban:
        quality:
          - type: text
            description: Must be a valid IBAN. Must not be empty.
```

SQL

An individual SQL query that returns a single number that can be compared with a threshold. The SQL query must be in the SQL dialect of the provided server.

Note: Establish a secure development process and use read-only connections, as the misuse of SQL queries can lead to SQL injection attacks.

Field	Type	Description
type	string	sql
description	string	A plain text describing the quality of the data.
query	string	A SQL query that returns a single number to compare with the threshold.
dialect	string	The SQL dialect that is used for the query. Should be compatible to the server type. Examples: postgres , spark , bigquery , snowflake , duckdb , ...
mustBe	integer	The threshold to check the return value of the query
mustNotBe	integer	The threshold to check the return value of the query

Field	Type	Description
mustBeGreaterThan	integer	The threshold to check the return value of the query
mustBeGreaterThanOrEqualTo	integer	The threshold to check the return value of the query
mustBeLessThan	integer	The threshold to check the return value of the query
mustBeLessThanOrEqualTo	integer	The threshold to check the return value of the query
mustBeBetween	array of two integers	The threshold to check the return value of the query. Boundaries are inclusive.
mustNotBeBetween	array of two integers	The threshold to check the return value of the query. Boundaries are inclusive.

In the query the following placeholders can be used:

Placeholder	Description
{model}	The name of the model that is checked.
{table}	Alias for {model} .
{field}	The name of the field that is checked (only if the quality is defined on field-level).
{column}	Alias for {field} .

Example:

```
models:
  orders:
    quality:
      - type: sql
        description: The maximum duration between two orders must be less that 360
        query: |
          SELECT MAX(duration) AS max_duration
          FROM (
            SELECT EXTRACT(EPOCH FROM (order_timestamp - LAG(order_timestamp) OVER
              FROM {model})
```

)
`mustBeLessThan: 3600`

SQL queries allow powerful checks for custom business logic. A SQL query should run not longer than 10 minutes.

Custom

You can define custom quality attributes that are specific to a data quality engine.

Custom (Engine: Soda)

Soda has a number of predefined quality [checks](#) that can be referenced as quality attributes.

Soda checks can be applied on model and field level.

Note: Soda Data contract check reference is experimental and may change in the future. Currently only supported by Postgres, Snowflake, and Spark (Databricks)

Field	Type	Description
type	string	custom
description	string	Optional. A plain text describing the quality attribute in natural language.
engine	string	soda
implementation	object	A check type as defined in the Data contract check reference

See the [Data contract check reference](#) for all possible types and configuration values.

Example:

```
models:
  my_table:
    fields:
      order_id:
        type: string
        quality:
          - type: custom
            description: This is a check on field level
            engine: soda
            implementation:
              type: no_duplicate_values
      carrier:
        type: string
      shipment_number:
        type: string
```

```

quality:
  - type: custom
    description: This is a check on model level
    engine: soda
    implementation:
      type: duplicate_percent
      columns:
        - carrier
        - shipment_number
      must_be_less_than: 1.0
  - type: custom
    description: This is a check on model level
    engine: soda
    implementation:
      type: row_count
      must_be_greater_than: 500000

```

Custom (Engine: Great Expectations)

Quality attributes defined as Great Expectations [Expectation](#).

Expectations are applied on model level.

Field	Type	Description
description	string	Optional. A plain text describing the quality attribute in natural language.
engine	string	great-expectations
implementation	object	An expectation type as listed in Expectation as YAML.

Example:

```

models:
  my_table:
    quality:
      - type: custom
        engine: great-expectations
        implementation:
          expectation_type: expect_table_row_count_to_be_between
          kwargs:
            min_value: 10000
            max_value: 50000
          meta:
            notes: "This expectation is crucial to avoid processing datasets that
      - type: custom
        engine: great-expectations
        description: "Check that passenger_count values are between 1 and 6."
        implementation:

```

```

expectation_type: expect_column_values_to_be_between
kwargs:
  column: passenger_count
  max_value: 6
  min_value: 1
  mostly: 1.0
  strict_max: false
  strict_min: false
meta:
  tags:
    - business-critical
    - range_check

```

Lineage Object

Field level lineage provides optional fine-grained information where the data comes from and how it was transformed.

The lineage object is based on the OpenLineage [Column Level Lineage Dataset Facet](#) to describe the input fields.

Field	Type	Description
inputFields	Array of InputField Object	The input fields refer to specific fields, columns, or data points from source systems or other data contracts that feed into a particular transformation, calculation, or final result.

InputField Object

Field	Type	Description
namespace	string	The input dataset namespace, such as the name of the source system or the domain of another data contract. Examples: <code>com.example.crm</code> , <code>checkout</code> , <code>snowflake://{account name}</code> . More on namespace
name	string	The input dataset name, such as a reference to a data contract, a fully qualified table name, a Kafka topic.
field	string	The input field name, such as the field in an upstream data contract, a table column or a JSON Path.

Field	Type	Description
transformations	Array of Transformation Object	Optional. This describes how the input field data was used to generate the final result.

Transformation Object

Field	Type	Description
type	string	Indicates how direct is the relationship e.g. in query. Allows values are: DIRECT and INDIRECT .
subtype	string	Optional. Contains more specific information about the transformation. Allowed values for type DIRECT : IDENTITY , TRANSFORMATION , AGGREGATION . Allowed values for type INDIRECT : JOIN , GROUP_BY , FILTER , SORT , WINDOW , CONDITIONAL .
description	string	Optional. A string representation of the transformation applied.
masking	boolean	Optional. Boolean value indicating if the input value was obfuscated during the transformation.

Example:

```
models:
  orders:
    fields:
      order_id:
        type: string
        lineage:
          inputFields:
            - namespace: com.example.service.checkout
              name: checkout_db.orders
              field: order_id
              transformations:
                - type: DIRECT
                  subtype: IDENTITY
                  description: The order ID from the checkout order
            - namespace: com.example.service.checkout
              name: checkout_db.orders
              field: order_timestamp
                - type: INDIRECT
                  subtype: SORT
          customer_email_address_hash:
            type: string
            lineage:
```



```
inputFields:
  - namespace: com.example.service.checkout
    name: checkout_db.orders
    field: email_address
    transformations:
      - type: DIRECT
        subtype: Transformation
        description: The email address from the checkout order, hashed w
        masking: true
```

Config Object

The config field can be used to set additional metadata that may be used by tools, e.g. to define a namespace for code generation, specify physical data types, toggle tests, etc.

A config field can be added with any name. The value can be null, a primitive, an array or an object.

For developer experience, a list of well-known field names is maintained here, as these fields are used in the Data Contract CLI:

Field	Type	Description
avroNamespace	string	(Only on model level) The namespace to use when importing and exporting the data model from / to Apache Avro.
avroType	string	(Only on field level) Specify the field type to use when exporting the data model to Apache Avro.
avroLogicalType	string	(Only on field level) Specify the logical field type to use when exporting the data model to Apache Avro.
bigqueryType	string	(Only on field level) Specify the physical column type that is used in a BigQuery table, e.g., NUMERIC(5, 2)
snowflakeType	string	(Only on field level) Specify the physical column type that is used in a Snowflake table, e.g, TIMESTAMP_LTZ
redshiftType	string	(Only on field level) Specify the physical column type that is used in a Redshift table, e.g, SMALLINT
sqlserverType	string	(Only on field level) Specify the physical column type that is used in a SQL Server table, e.g, DATETIME2
databricksType	string	(Only on field level) Specify the physical column type that is used in a Databricks table

Field	Type	Description
glueType	string	(Only on field level) Specify the physical column type that is used in a AWS Glue Data Catalog table

This object *MAY* be extended with [Specification Extensions](#).

Example:

```
models:
  orders:
    config:
      avroNamespace: "my.namespace"
    fields:
      my_field_1:
        description: Example for AVRO with Timestamp (millisecond precision)
        type: timestamp
        config:
          avroType: long
          avroLogicalType: timestamp-millis
          snowflakeType: timestamp_tz
```

Data Types

The following data types are supported for model fields and definitions:

- Unicode character sequence: `string` , `text` , `varchar`
- Any numeric type, either integers or floating point numbers: `number` , `decimal` , `numeric`
- 32-bit signed integer: `int` , `integer`
- 64-bit signed integer: `long` , `bigint`
- Single precision (32-bit) IEEE 754 floating-point number: `float`
- Double precision (64-bit) IEEE 754 floating-point number: `double`
- Binary value: `boolean`
- Timestamp with timezone: `timestamp` , `timestamp_tz`
- Timestamp with no timezone: `timestamp_ntz`
- Date with no time information: `date`
- Array: `array`
- Map: `map` (may not be supported by some server types)
- Sequence of 8-bit unsigned bytes: `bytes`
- Complex type: `object` , `record` , `struct`
- No value: `null`

Specification Extensions

While the Data Contract Specification tries to accommodate most use cases, additional data can be added to extend the specification at certain points.

A custom field can be added with any name. The value can be null, a primitive, an array or an object.

Tooling

- [Data Contract CLI](#) is an open-source CLI tool to help you create, develop, and maintain your data contracts.
- [Data Contract Manager](#) is a commercial tool to manage data contracts. It includes a data contract catalog, a Web-Editor, and a request and approval workflow to automate access to data products for a full enterprise data marketplace.
- [Data Contract GPT](#) is a custom GPT that can help you write data contracts.
- [Data Contract Editor](#) is an open-source editor for Data Contracts, including a live html preview.

Code Completion

The [JSON Schema](#) of the current data contract specification is registered in [Schema Store](#), which brings code completion and syntax checks for all major IDEs. IntelliJ comes with a built-in YAML plugin which will show you autocompletions. For VS Code we recommend to install the [YAML](#) plugin. No additional configuration is required.

Autocompletion is then enabled for files following these patterns:

```
datacontract.yaml
datacontract.yml
*-datacontract.yaml
*-datacontract.yml
*.datacontract.yaml
*.datacontract.yml
datacontract-*.yaml
datacontract-*.yml
**/datacontract/*.yaml
**/datacontract/*.yml
**/datacontracts/*.yaml
**/datacontracts/*.yml
```

Authors

The Data Contract Specification was originally created by [Jochen Christ](#) and [Dr. Simon Harrer](#), and is currently maintained by them.

Contributing

Contributions are welcome! Please open an issue or a pull request.

License

[MIT License](#)

This site is open source. [Improve this page.](#)



[Legal Notice](#) | [Privacy](#)